

Figure 1 depicts the architectural details of Apache Airavata. From this diagram one can see how end-users are able to interact with the Apache Airavata services (API Server, Workflow Interpreter, Orchestrator, Messaging System and etc.) via Science Gateways. The Apache Software foundation provides a detailed description of the architectural diagram shown in figure 1 [2].

## API

As introduced in section 4 and shown with multiple real-world examples in section 5, researchers leverage Science Gateways to interact with Apache Airavata. In order to promote simplicity, Apache Airavata's application programming interface (API) is intentionally obscured from these end-users. Therefore, the Airavata API is generally intended for Science Gateway developers who are specifically interested in using Apache Airavata as a middleware service between a user interface and one or more compute resources, as shown in figure 2. Airavata API is written using apache thrift [2]. This allows Science Gateway developers to use the programming language of their choice (e.g. Java, PHP, JavaScript, C++, etc.). The Apache Software Foundation provides an in-depth overview of the API for those interested in learning the details of this service [3].

## Shell Access

As section 1 thoroughly explained, Apache Airavata's purpose is to simplify the typical command line driven process of composing, managing, executing, and monitoring large-scale scientific applications on powerful distributed computing resources. Therefore, end-users of Apache Airavata should not rely on shell access. Instead, section 2.3 explains that end-users interact with Apache Airavata through graphical interfaces (i.e. Science Gateways). Shell access is contained to the Science Gateway developer level of the technological ecosystem, shown in figure 2.

## Graphical Interface

Apache Airavata leverages a Graphical Workflow Composer, known as XBay, which helps "create workflows, submit and manage multiple applications ... [and] also has a web-based interface ... where users can ... register, run and monitor applications " [2]. Furthermore, the Apache Airavata's thrift-based API, introduced and discussed in section 2.1, allows developers to create their own desktop and web interfaces using Airavata as the technological foundation.

## Science Gateways

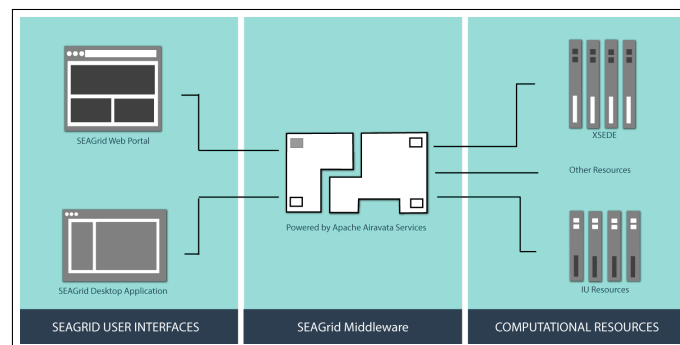
Science Gateways are the resulting reward of leveraging Apache Airavata. As explained in 4, Science Gateways are an instrumental piece of the Apache Airavata ecosystem that allow end users to compose, manage, execute, and monitor scientific research workflows on distributed (and potentially complex) compute resources. Science Gateways simplify workflows and allow researchers to focus their expensive time and effort on science. Furthermore, Science Gateways promote reproducibility, which is an important piece of scientific research as well as publication. In other words, it is beneficial to researchers to have the ability to easily reproduce results from an experiment in the past and Science Gateways typically ensure this functionality. Section 5 describes many of the prominent and currently available Science Gateways built on top of the Apache Airavata software framework.

## LICENSING

Apache Airavata is open source software [4]. Therefore, anyone can download, install, modify and improve this software framework using the popular "fork and pull" model. As part of an introductory tutorial, one can create an Airavata Test Drive account, which is free as well. Note, this free Airavata Test Drive

account will need to be approved by an Airavata Test Drive Administrator before one can use the system. Once this approval occurs, one can simply create and run "experiments" (i.e. Gaussian, Amber, Trinity and etc.) via a Science Gateway on local clusters, supercomputers, computational grids and computing clouds.

## ECOSYSTEM



**Fig. 2.** The image above depicts an example of the technological ecosystem developed around Airavata [5].

Figure 2 depicts the ecosystem developed around Apache Airavata for SEAGrid. Apache Airavata plays an important role as the middleware between the end-users (Science Gateways) and the compute resources. As discussed above, the placement of Apache Airavata in between the compute resources and the science gateways allows for the abstraction of the complexity of composing, running, executing and managing applications and workflows.

## USE CASES

Ultrascan [6], SEAGrid [7] and GenApp [8] are instances of Science Gateways that leverage Apache Airavata to perform computations [1]. Each of these services is in place to simply bridge the gap between scientific applications on large-scale compute resources and domain specialists. In other words, these Science Gateways are examples of Apache Airavata enabling science.

## SEAGrid

The Apache Airavata services allow Science Gateways such as SEAGrid to simplify the use of "scientific applications deployed across a wide range of supercomputers, campus clusters and computing cloud" [7]. SEAGrids bridges the gap between domain specialists and scientific applications on large-scale compute resources using both a desktop client and web application. SEAGrid currently promotes research in Computational Chemistry (e.g. Gaussian, Gamess, etc.), Molecular Dynamics (Lammps, Amber, NAMD and etc.), Structural Mechanics (e.g. Abaqus and etc.), Fluid Dynamics (e.g. Nek5000, OpenFOAM and etc.) and much more. SEAGrid abstracts away the fine-grained details of running such scientific applications on large-scale compute resources and therefore allows the domain specialists to focus on the fine-grained details of their scientific research. Additionally, SEAGrid enables scientists to create model inputs, visualizations of outputs and archives for simulation data" [7].

## Use Cases for Big Data

The One-Degree Imager (ODI) “is a gigapixel mosaic camera ... built by [the] WIYN Observatory with a pixel scale of 0.1 arcseconds for the 3.5-meter telescope” and is the newest instrument at the WIYN 3.5m Observatory in Sells, AZ [9, 10]. Similarly to the SEAGrid Science Gateway, the Apache Airavata software framework is at the foundation of ODI’s Pipeline, Portal and Archive (PPA) system which “execute[s] the NOAO High Performance Pipeline System (NHPPS) pipelines on XSEDE resources” [9]. The large amount of data generated by the ODI demonstrates that the ODI-PPA Science Gateway and therefore the Apache Airavata software framework can handle big data software projects.

Don’t use And as a sentence starter, a special recommendation from professor.(Last sentence in the Educational material paragraph).

## EDUCATIONAL MATERIAL

There are multiple ways to find out more information about the Apache Airavata software framework [1]. In order to cater to varying audiences and learning styles there is online documentation [2], a course provided at Indiana University, Bloomington [11] as well as online tutorials [12]. The online documentation is entirely sufficient for motivated users to teach themselves. Historically, there has been a basic Airavata course at Indiana University offered during the Fall semester and a more advanced version of the course is offered during the Spring semester. And finally, the online tutorials provide everything from quick-start to extended tutorials and everything in between.

## CONCLUSION

Apache Airavata appeals to a wide range of scientific researchers since the technology allows researchers to focus their time, effort and grant money on the science rather than the details of the computing. Apache Airavata has enabled researchers to compose, manage, execute and monitor workflows on large-scale systems with the click of a button. Gateways such as Ultrascan, SEAGrid and GenApp are clear and defined examples of how Apache Airavata has been leveraged to improve and optimize scientific research workflows. As compute resources get more complicated and/or distributed over time, the Apache Airavata software framework will continue to promote the ease of use with Science Gateways.

## ACKNOWLEDGEMENTS

The authors would like to thank the School of Informatics and Computing for providing the Big Data Software and Projects (INFO-I524) course [13]. This paper would not have been possible without the technical support & edification from Gregor von Laszewski and his distinguished colleagues.

## AUTHOR BIOGRAPHIES



**Scott McClary** received his BSc (Computer Science) and Minor (Mathematics) in May 2016 from Indiana University and will receive his MSc (Computer Science) in May 2017 from Indiana University. His research interests are within scientific application performance analysis on large-scale HPC systems. He will begin working as a Software Engineer with General Electric Digital in San Ramon, CA in July 2017.

## WORK BREAKDOWN

The work on this project was distributed as follows between the authors:

**Scott McClary.** He completed all of the work for this paper including researching and testing Apache Airavata as well as composing this technology paper.

## REFERENCES

- [1] Apache Software Foundation, “Apache Airavata,” Web Page, 2016, accessed: 2017-2-22. [Online]. Available: <http://airavata.apache.org>
- [2] —, “Apache Airavata Overview - Apache Airavata - Apache Software Foundation,” Web Page, February 2016, accessed: 2017-2-22. [Online]. Available: <https://cwiki.apache.org/confluence/display/AIRAVATA/Apache+Airavata+Overview>
- [3] —, “Airavata API Overview - Apache Airavata - Apache Software Foundation,” Web Page, May 2014, accessed: 2017-2-22. [Online]. Available: <https://cwiki.apache.org/confluence/display/AIRAVATA/Airavata+API+Overview>
- [4] —, “Airavata Source Code & Developers Guide,” Web Page, 2014, accessed: 2017-2-22. [Online]. Available: <https://airavata.apache.org/development/source.html>
- [5] Indiana University, Apache Airavata, XSEDE, and NSF, “SEAGrid Portal,” Web Page, accessed: 2017-2-22. [Online]. Available: <https://seagrid.org/themes/seagrid/assets/img/workflow.png>
- [6] UltraScan Project, “UltraScan Analysis Software,” Web Page, April 2015, accessed: 2017-2-22. [Online]. Available: <http://ultrascan.uthscsa.edu>
- [7] Indiana University, Apache Airavata, XSEDE, and NSF, “SEAGrid Portal,” Web Page, accessed: 2017-2-22. [Online]. Available: <https://seagrid.org>
- [8] Apache Software Foundation, “GenApp - Apache Airavata - Apache Software Foundation,” Web Page, August 2014, accessed: 2017-2-22. [Online]. Available: <https://cwiki.apache.org/confluence/display/AIRAVATA/GenApp>
- [9] —, “ODI - Apache Airavata - Apache Software Foundation,” Web Page, January 2017, accessed: 2017-2-22. [Online]. Available: <https://cwiki.apache.org/confluence/display/AIRAVATA/ODI>
- [10] WIYN Observatory, “WIYN 3.5m Observatory,” Web Page, February 2017, accessed: 2017-2-22. [Online]. Available: <https://www.noao.edu/wiyn/ODI/>
- [11] M. Pierce and S. Maru, “I590 Course Structure,” Web Page, accessed: 2017-2-22. [Online]. Available: <http://courses.airavata.org/spring2016/public/slides/I590CourseStructure.pdf>
- [12] Apache Software Foundation, “Airavata Quick-Start Tutorials - Apache Airavata - Apache Software Foundation,” Web Page, July 2016, accessed: 2017-2-22. [Online]. Available: <https://cwiki.apache.org/confluence/display/AIRAVATA/Airavata+Quick-Start+Tutorials>
- [13] Gregor von Laszewski and Badi Abdul-Wahid, “Big Data Classes,” Web Page, Indiana University, Jan. 2017. [Online]. Available: <https://cloudmesh.github.io/classes/>