

Xen: A bare metal hypervisor

PIYUSH RAI¹

¹School of Informatics and Computing, Bloomington, IN 47408, U.S.A.

*Corresponding authors: piyurai@iu.edu

+HID - S17-IO-3014

April 7, 2017

Xen is an open-source baremetal hypervisor. This paper explores it's overview, architecture and contains a brief note on its alternatives.

© 2017 <https://creativecommons.org/licenses/>. The authors verify that the text is not plagiarized.

Keywords: Baremetal hypervisor, Xen

<https://github.com/piyurai/sp17-i524/tree/master/paper1/S17-IO-3014/report.pdf>

INTRODUCTION

To provide an isolated environment to different applications in terms of memory, process scheduling and disk access such that one process does not impact the other, used to be a major challenge for system administration [1]. Also, different applications may have different OS requirements. Running an application within its own OS environment should protect it from other malicious applications because of the guaranteed memory and disk allocation to each guest VM (virtual machine).

Xen is a baremetal hypervisor and is available as open-source [2]. It's based on microkernel design with the advantage of small memory footprint and has a limited guest interface. It's responsibility is to manage CPU and memory, and to handle interrupts. Virtual machines are deployed in the guest domain called DomU. A special virtual machine is deployed in the control domain called Domain 0. It contains hardware drivers and the toolstack to control the VMs.

ARCHITECTURAL OVERVIEW

Xen has a small memory footprint and provides limited interface to guest VMs. It runs at the highest CPU privilege level while the guests are deployed in DomU domain [2]. The hypervisor is the first process to be started after bootloader. It's primary responsibility is to manage CPU, memory and interrupts and does not contain the drivers for I/O functions such as networking and storage. The main device driver for a system can be run inside of a virtual machine. It protects the rest of the system from events such as driver crashes as the VM containing the driver can be rebooted independently. The control stack is generally deployed on a Linux VM running in domain 0. It is the first VM to be deployed by the system and can access the hardware directly. It handles system's I/O functions and provides interface to control the system. It's responsible for creating and configuring VMs, allocating resources and monitoring them and terminating the

VMs when they are no longer required. The control stack interface can also be driven through a cloud orchestration stack such as OpenStack. Figure 1 shows the architectural overview of Xen.

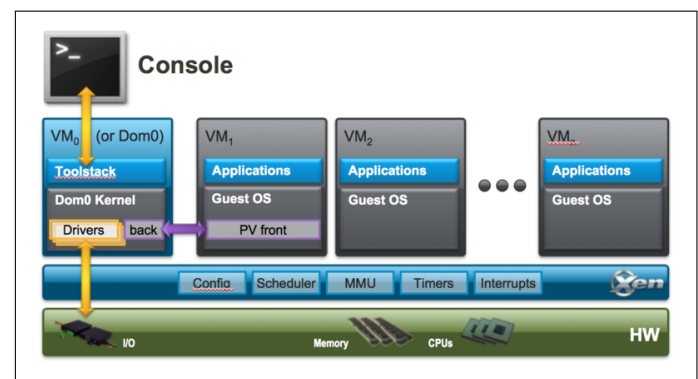


Fig. 1. Xen architecture overview [3].

VIRTUALIZATION MODES

The guest virtual machines are deployed in domain DomU which has no privilege to access the hardware. The guest VMs can either be deployed in Para virtualization (PV) mode or Hardware-assisted virtualization mode (HVM). It's also possible for a VM to use PV drivers inside of a HVM mode to improve its performance. The two types of modes can be deployed simultaneously on a single hypervisor. PV guests work without virtualization support from the hardware, and require their kernel and drivers to be PV-enabled. However, the application binaries do not need any modification and run as in their native environments.

HVM uses virtualization support available on the host CPU such as Intel VT or AMD-V hardware extensions. QEMU, an

open source system emulator, is used to emulate the remaining PC hardware. No changes to kernel are required for fully virtualized guests. HVM guests can also use PVHVM drivers i.e. special paravirtual device drivers optimized for HVM environments to boost up the performance.

PVH is a recently developed new virtualization mode, introduced in Xen Project 4.4 where PV guests uses PV drivers for boot and I/O and hardware extensions otherwise. It is meant to reduce the number of interfaces available to guest thereby reducing the attack surface. It is expected to simplify the Xen architecture and is the recommended virtualization mode to be used with newer Xen releases [4].

RESOURCES

[5] presents an overview of the software architecture and host the links to documentation and release notes for the different releases. [6] explains the hypervisor internals and it's important features along with their interfaces.

ALTERNATIVES

The early days of Xen development came across different limitations such as modification to the guest kernels and limited Linux support [7]. Into the later years of the development of Xen, it remained based on the older version of Linux kernel as the developers waited for more recent kernels with support for modern hardware [8]. This allowed KVM to overtake as preferred choice for virtualization. KVM overcame the problems faced by Xen by using the virtualization support in the newly available hardware and reusing the components from the Linux kernel. This gave it the advantage of automatically benefiting from the developments in the mainline kernel. Its incorporation into the Linux version 2.6.20 and contribution from companies like AMD and Intel increased KVM's adaptability among the Linux distributors.

KVM requires hardware assisted virtualization with new releases containing paravirtualization support for certain devices [9]. Xen was originally developed to overcome the problem of limited support for hardware assisted virtualization [4] by using Paravirtualization. Xen is a layer between the kernel and hardware whereas KVM turns the kernel itself into the hypervisor. Xen has it's own memory and power management system designed specifically for VMs where as KVM needs to have support for processes as well.

[10] lists the observations made during comparisons of different virtualization technologies in 2013 using DTrace. It analyzes the stack trace for different system functionalities such as network and I/O calls stack trace along with the relevant overhead for Zones, Xen and KVM. KVM took fewer steps than Xen in terms of I/O path. Zones, with Linux Containers being its equivalent, turned out to be the preferred choice for virtualization as it had the lowest overhead for not having to run an entire OS for each application.

CURRENT STATUS AND THE FUTURE

Xen is actively maintained by Linux Foundation under the trademark "XEN Project". Some of the features included in the latest releases include "Reboot-free Live Patching" (to enable application of security patches without rebooting the system) and KCONFIG support (compilation support to create a lighter version for requirements such as embedded systems) [11]. The overheads in terms of passing data across several layers is being

minimized by use of shared memory transports and buffers [10]. Xen Project also strives to minimize the footprint of the change to Linux kernel. The codebase contains less than 150,000 lines of code [2]. Xen Project hosts a listing of vendors that uses Xen Project software to allow them to advertise their software and services. Xen is used by more than 10 million users [12]. Some of the major cloud hosting service providers using Xen include Amazon and Rackspace. Hypervisors are predicted to become a core aspect of cloud based solutions. Containers are growing in popularity and there's an increasing interest in combination of isolation aspects of hypervisor and container runtime environments [13]. [11] talks about the ongoing work within Xen community towards the development of hypervisor-based containers.

REFERENCES

- [1] P. Barham, B. Dragovic, K. Fraser, S. Hand, T. Harris, A. Ho, R. Neugebauer, I. Pratt, and A. Warfield, "Xen and the art of virtualization," in *Proceedings of the nineteenth ACM symposium on Operating systems principles*, ser. SOSP '03. New York, NY, USA: ACM, 2003, pp. 164–177. [Online]. Available: <http://dl.acm.org/citation.cfm?id=945462>
- [2] "Xen Project Software Overview," Web Page, Mar. 2016. [Online]. Available: https://wiki.xenproject.org/wiki/Xen_Project_Software_Overview
- [3] "Xen Architecture Diagram," Web Page, Apr. 2015. [Online]. Available: https://wiki.xenproject.org/wiki/File:Xen_Arch_Diagram.png
- [4] David Chisnall, "Xen PVH," Web Page, Jun. 2014. [Online]. Available: <http://www.informit.com/articles/article.aspx?p=2233978>
- [5] "Xen Project wiki," Web Page, Dec. 2016. [Online]. Available: https://wiki.xenproject.org/wiki/Main_Page
- [6] D. Chisnall, *The Definitive Guide to the Xen Hypervisor*, 1st ed., ser. Open Source Software Development Series, A. Robbins, Ed. Prentice Hall, nov 2007, no. ISBN-10: 0-13-358249-3.
- [7] Amit Shah, "Ten years of KVM," Web Page, Nov. 2016. [Online]. Available: <https://lwn.net/Articles/705160/>
- [8] Thorsten Leemhuis, "Rise of KVM," Web Page, Jun. 2011. [Online]. Available: <http://www.h-online.com/open/features/Xen-lets-KVM-overtake-1262171.html>
- [9] "KVM - Wikipedia," Web Page, Feb. 2017. [Online]. Available: https://en.wikipedia.org/wiki/Kernel-based_Virtual_Machine
- [10] Brendan Gregg, "Virtualization Performance: Zones, KVM, Xen," Web Page, Jan. 2013. [Online]. Available: <http://dtrace.org/blogs/brendan/2013/01/11/virtualization-performance-zones-kvm-xen/>
- [11] "Xen Project 4.7 Feature List," Web Page, Jun. 2016. [Online]. Available: https://wiki.xenproject.org/wiki/Xen_Project_4.7_Feature_List
- [12] Stefano Stabellini, James Bulpin, "Hypervisor in 2017," Web Page, Dec. 2016. [Online]. Available: goo.gl/S3X1oQ
- [13] Stefano Stabellini, "Hypervisor-Based Containers," Web Page, Dec. 2016. [Online]. Available: <https://thenewstack.io/hypervisors-container-era/>