

Apache Lucene

ROY CHOUDHURY, SABYASACHI^{1,*}, +

¹School of Informatics and Computing, Bloomington, IN 47408, U.S.A.

*Corresponding authors: sabyroyc@indiana.edu

+HID - S17-IO-3015

project-000, March 15, 2017

This paper gives an overview of Apache Lucene. We will go through the basic architecture and functionality of the library. We will see the advantages and downfalls of Lucene and conclude on its implementation.

© 2017 <https://creativecommons.org/licenses/>. The authors verify that the text is not plagiarized.

Keywords: Search Engine Library, Lucene

<https://github.com/sabyasachi087/sp17-i524/tree/master/paper1/S17-IO-3015/report.pdf>

This review document is provided for you to achieve your best. We have listed a number of obvious opportunities for improvement. When improving it, please keep this copy untouched and instead focus on improving report.tex. The review does not include all possible improvement suggestions and for each comment you may want to check if it applies elsewhere in the document.

Abstract: Please redo. The abstract needs to be about the technology, not about the paper itself. You should say briefly what the most important things about Lucene are. Then maybe in the last sentence, you can tell the reader what the paper will cover. This is the part of the paper that the most people will read, and will use it to decide whether to read on. As it stands, someone reading your abstract will still have no idea what Lucene is if they didn't before.

Please, format lines to 80 characters in the source to make the LaTeX easier to read.

You need to expand your paper. Currently, it's closer to one page, excluding diagram and references. It needs to be two pages.

Overall, your paper needs to be more detailed when describing specific functions. In several places you use very general sentences that don't provide information, and someone would have to already know about search engines and Lucene to evaluate them. Please see the comments below for details.

Assessment: Revisions required. Please address the review comments by end of March.

INTRODUCTION

Apache Lucene is a search library that enables search facility to any application. Although it was initially written in Java but has been

Grammar

ported to many languages chiefly C-Sharp, c/c++

Grammar

, Python etc. It is an active open source project and under Apache License. Its latest release is 6.4.1.

The version is out of scope for the paper. The paper should be able to stand more generally, independent of when it was written. If you think it's really important to mention the version (e.g. there are significant changes after a given version, you can state that in a footnote. That way you won't break up the flow of the paper with information that most readers won't care about.

It is important that we must understand that

The beginning of this sentence is unnecessary. Try to keep your sentences shorter and to the point; it will make reading the paper easier.

Lucene is just a search library and cannot handle other related stuffs like crawling, document filtering, administration etc.

CONCEPTS

To understand the purpose of Lucene we have to be familiar with two terms, one is "Information Overload" and "Information Retrieval"[1]

Leave a space between the referece and the previous word.

When you introduce terms, if you want to draw attention to them, it's better to use the LaTeX `\emph{}` command, rather than quotes.

Like before, you should try to get straight to the point. This first sentence doesn't help the reader understand anything about Lucene, IR or information overload, so it can be skipped. You can start by describing information overload like you do later in this paragraph.

The term information overload means, difficulty that one can have in making decisions, because of the presence of too much of information. In another

Term

words, "Information Overload"

You've already introduced the term so no need for quotes or special formatting anymore.

can occur if the rate of feed/input into a system exceeds its processing capabilities. Imagine the situation of current world. We are living in the world where data has reached volume of zeta bytes. To extract some insight, first step is to collect all related data. This is known as information retrieval (IR). IR is the task of collecting relevant information from collection of data resources scattered across devices. Now with the above understanding we can say Lucene is a scalable Information Retrieval library.

Lucene and Search Engine

Google has set some base expectations for search engines, which if not available can cause dissatisfaction for the users. For example spell checker and response time of ≤ 1 second. Lucene should be able to met those expectations or else its

Term

Know when to use "it is" versus "its"

not worth.

Grammar

But

Avoid starting a sentence with "and", "but", "so".

Lucene is not a full fledged search engine rather its a tool kit to achieve so.

Grammar

So

Term

lets jot down the steps for any search engine

"jot down" is too conversational for a paper like this.

Module : Raw Content -> Gather Data -> Analyze -> Index -> Query Support UseCase : User query -> Search Engine -> Build Query -> Extract Information from indexed data

This can be expanded significantly. It basically describes what a search engine does, but provides almost no detail. If someone reading this doesn't already know how search engines work, they won't be helpful.

ARCHITECTURE

Figure 1 explains a basic architecture of Lucene

Lucene library is compact and does not have any external dependencies. But

Term

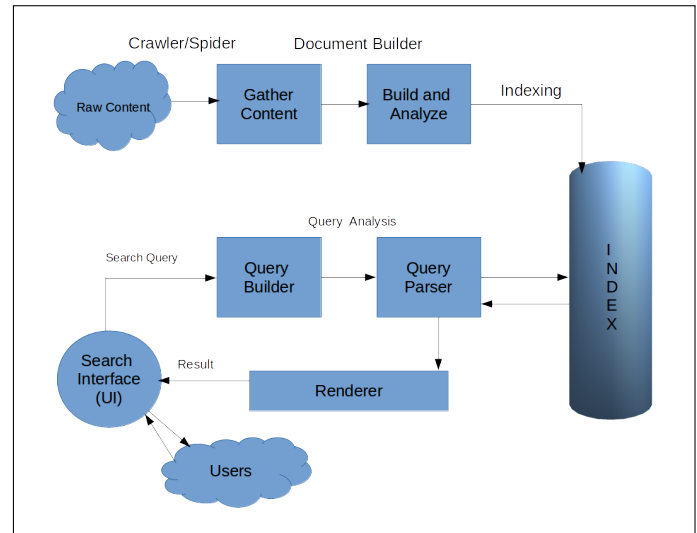


Fig. 1. Apache Lucene Architecture.

Did you create this diagram? If not, you need to add a reference to where you got it from.

it can be plugged

Term

with other libraries for building a search engine.

What other libraries? You've already said Lucene doesn't provide everything needed for a search engine, so you need to explain exactly what it provides, and what other technologies someone working on a search engine will need. Imagine you are a user who wants to set up a search engine for their data. Other than Lucene, what else will they need?

Fig 1[2] shows process flow of Lucene.

Grammar

On a high level, data collected from different sources are analyzed and converted into smaller chunks. These are known as documents. Documents are text entries and from these text entries Lucene performs indexing and store it in local disk for future reference.

What is the purpose of indexing?

The next step is to handle the search query from users. Lucene has a query parser to understand the query and search the index for the correct or relevant match. If found, returns the document back. We will elaborate the flow beneath.

LUCENE COMPONENTS

Document Analysis and Indexing

Data or contents which are available in different format and location needs to be gathered. This process is typically done by a crawler/spider. Core Lucene does not have these capabilities and can be considered to be a pre-requisite for Lucene to implement.

It's important to mention that Lucene can be made to work on any document collection or database. You don't need to crawl anything from the Internet.

Two of the crawler that are build on Lucene are Solr and Nutch.

These need to be added as references in references.bib, not simply as links.

Once the data is collected , document has to be constructed from the contents. The design of constructing documents has to be decided by the user and implemented within Lucene. Lucene provides an API for building documents but logic has to be provided by the implementation layer. Lucene also does not provide any API for document filtering.

What is document filtering? How does it fit into the basic process of collecting the data, preprocessing it for Lucene, and Lucene indexing it?

But yet again we have Tika , which is built on Lucene, can be used for this purpose. But we cannot index the document yet.

This and the next sentence say the same thing. Redundancy.

We cannot index the raw content within the document directly. Before that , we need to break the content into smaller chunks known as tokens. Each token is map to a "word".

Not clear what this sentence says. Is a token a word? Why is "word" quoted?

Analysis includes handling compound words, spell check, typo correction injection of synonyms, etc. Lucene has built in support of list of analyzers

Grammar

Where do these analyzers fit in the Lucene architecture?

which gives a fine grain control over analysis. Once tokenization is done , now

Term

its

"it is"

time for indexing. Lucene takes care all of the need to cater this step. An API has been provide for this purpose, but has to be implemented carefully as the searching will solely depend upon how well the indexing has been done.

You need to explain what exactly indexing is. This portion is a bit too general and is not clear if you don't already know what indexing is.

Searching

It is a look up process

Need a full sentence, don't refer to the section title. "Searching is the process of looking up..."

within the index to extract the most relevant (matching) documents. It is based on two matrices

Do you mean metrics?

i.e.

Term

Precision and Recall. Recall measures how well the system finds the relevant documents and Precision measures the filtering out the irrelevant one.

Not a good explanation of precision and recall. Recall measures the fraction of relevant documents that are retrieved. Precision measures the fraction of retrieved documents that are relevant.

Lucene offers benchmarking technique for measuring these matrices

Term

. User Interface (UI) is equally important for a search application as that is what the end user is going to use. Lucene does not provide any UI support. When user inputs the search query the first step is to build the query. User inputs are human readable and need further processing before it can be used within the application. Lucene provides a powerful parser for this job known as QueryParser. Even it is default state

Grammar

it does the job pretty clean

"pretty clean" is too vague.

but often needs extension as per some advance requirements. Finally hitting the search query to index.

This is not a complete sentence.

Almost everything about this is catered by Lucene and also provide option for extension.

"Almost everything about this is catered by Lucene" is too general; it doesn't provide any useful information

It finds the result and returns all relevant document objects. Its

"it is"

the responsibility of the UI to render the results correctly.

ADVANCE USAGES

Term

Lucene has some advance feature for administration and analytics. For example Lucene allows to configure the RAM buffer size , re-indexing ,commit and purge scheduler.

You need to provide a little more detail what these things are.

It allows some fault tolerance mechanism in case a newly added document failed to index. Related to analytics it also provides some meta information regarding the search queries it receives and the results it renders. For example which kind of query are run , query hitting lowest relevance , query having no results and so on. One big problem within the list of advance usages

Term

that Lucene does not support is "Scaling"

Don't use quotes.

. Scaling in terms of both through put and processing speed. In a clustered environment this is quiet

Term

an important part as data and resource all are distributed. But both Solr and Nutch provides data partitioning and sharding to achieve higher throughput if not speed. Elastic search

What is elastic search? You either need to explain briefly or provide a reference.

is another option which is based upon Lucene and provides distributed computing.

CONCLUSION

Lucene is the standard library for search applications. It can be compared with 'C' (language) of computing which is small and powerful but requires much more

Term

effort to build an entire application.

The comparison is a stretch, and subjective.

It must be remembered that Lucene is just a library which sits at the core of the functionality but it needs much more

"much more" to vague; avoid using adjectives like "much" in a paper like this.

than that to build an application. Elastic Search , Solr and Nutch which are based on Lucene are preferred tools in terms of building enterprise level search engines.

READING SOURCES

- Lucene In Action [2] provides a startup guide to learn, build and implement search engines based on Lucene
- Lucene at tutorial point [3] provides introduction to Lucene Library.

ACKNOWLEDGEMENTS

Thanking Prof. Gregor von Laszewski for his technical help and support.

REFERENCES

- [1] "Wikipedia," Web page. [Online]. Available: https://en.wikipedia.org/wiki/Information_retrieval
- [2] O. G. Erik Hatcher, *Lucene In Action*, 2nd ed. Manning, 2004.
- [3] "Tutorialpoint," Web page. [Online]. Available: <https://www.tutorialspoint.com/lucene/index.htm>