

Docker(Machine,Swarm)

SHREE GOVIND MISHRA¹

¹School of Informatics and Computing, Bloomington, IN 47408, U.S.A.

* Corresponding authors:shremish@indiana.edu and govindmisharaalld@gmail.com

project-000, February 27, 2017

Docker is a container based technology which helps in the packaging and shipping of applications quickly and across networks. Docker is being well accepted and appreciated in the Agile Software Development workforce and among the DevOps because of their attributes of Continuous Integration and testing. To create a cluster of Docker Engines into a single docker virtual engine is used Docker-Swarm which is tested for more than 50,000 containers. © 2017 <https://creativecommons.org/licenses/>. The authors verify that the text is not plagiarized.

Keywords: Cloud, I524, Big Data, Virtualisation, Docker, Docker Swarm, Docker Machine, Virtual Machines, VMs, Containers, Linux Containers, lightweight containers, LXC

<https://github.com/argetlam115/sp17-i524/blob/master/paper1/S17-IR-2021/report.pdf>

INTRODUCTION

[1] Docker is an open-source container based technology. It is an extension of Linux Containers (LXC): which is a unique kind of lightweight, application-centric virtualization that drastically reduces overhead and makes it easier to deploy software on servers. A container allows a developer to package up an application and all its part including its code, stack it runs on, dependencies it is associated with, system tools and everything the application requires to run within an isolated environment. This makes it easier for programmers and developers to run more apps on the same server and it even makes it easier to package and ship the apps very frequently. Docker has been able to popularize the container approach in part because it's improved the security and simplicity of container environments. Plus, interoperability is enhanced by its association with major companies – such as Google, Canonical, and Red Hat – on its open source element libcontainer.

HOW DOCKER DIFFERS FROM VMS

The virtualisation of application can be obtained with Hypervisors or Virtual Machine Manager (VMM) which makes it easy for applications to run in isolation with one another while sharing the same underlying hardware architecture. VM hypervisors, such as Hyper-V, KVM, and Xen, all are "based on emulating virtual hardware" which means they're fat in terms of system requirements. Instead of virtualizing hardware, containers rest on top of a single Linux instance. A full virtualized system gets its own set of resources allocated to it, and does a very minimal sharing of those resources. So, the developer gets more isolation but each instance of the VM also has many "Junk VM files", which is not useful to the developers as it gets "heavy". Since they

only require to keep the virtual machine image and it can be kept small. Thus, the smaller the image is the less we need to store and the less you need to send around the network which makes them fairly lightweight in comparison to the Virtual Machines. Thus, we can run even use 1000s of containers on a same host OS. [2] A full virtualized system usually takes time to start whereas the Docker container does take even less than seconds to up start and running with a lower overhead than the VMs. Containers are potentially much more efficient than VMs because they're able to share a single kernel and share application libraries. This can lead to substantially smaller RAM footprints even when compared to virtualisation systems that can make use of RAM overcommitment. Storage footprints can also be reduced where deployed containers share underlying image layers. IBM's Boden Russell has done

USES OF DOCKER

Docker for DevOps

Another major use of Docker is its use in the DevOps community. Docker is not there to replace other configuration management tools and instead can be incorporated with other configuration management tools like Chef, Puppet, Salt or Ansible. The other major benefit of using Docker, Dockerfiles, the registry and the whole Docker ecosystem is that the teams don't have to learn domain specific language as these are easier to learn than the domain specific Ecosystems. Though many a times Docker can be made to work with the other configuration management tools. [3] Docker can also eliminate the need for a development team to have the same versions of everything installed on their local machine. The repeatable nature of Docker images makes it easier for them to standardize their production code and configurations. Their work has led to the creation of Helios,

an application that manages Docker deployments across multiple servers and that alerts them when a server isn't running the correct version of a container.

Docker as Virtualized Sandbox

[3] Docker allows systems administrators and developers to build applications that can be run on any Linux distribution or hardware in a virtualized sandbox without making custom builds for all the different environments. Finally, it's easy to deploy Docker containers in a cloud scenario. You can easily integrate it with typical DevOps environments seamlessly (Ansible, Puppet, etc.) or use it as a standalone.

Docker for continuous integration

Docker can be used as git for continuous integration. Docker is similar as changes in the system can be tracked just like changes in the git. These collaboration features (docker push and docker pull) are one of the most disruptive parts of Docker. The fact that any Docker image can run on any machine running Docker is very much appreciated. But the Docker pull/push are the too new for the first time developers and ops guys have ever been able to easily collaborate quickly on building infrastructure together.[1] The app guys can share app containers with ops guys and the ops guys can share MySQL and PostgreSQL and Redis servers with app guys.

DOCKER MACHINE

[4] Docker Machine is a tool that lets you install Docker Engine on virtual hosts, and manage the hosts with docker-machine commands. You can use Machine to create Docker hosts on your local Mac or Windows box, on your company network, in your data center, or on cloud providers like AWS or Digital Ocean. Docker Machine is the only way to run Docker Engine on Mac or Windows previous to Docker v1.12 and starting with the Docker v1.12 we have Docker for Mac and Docker for Windows. Thus we can create a cluster of Docker hosts which is called a Swarm using Docker Machine.

DOCKER SWARM

[5] Docker Swarm provides native clustering capabilities to turn a group of Docker engines into a single, virtual Docker Engine. These help you scale up the applications as if these all are running on a single, huge computer. It does so by providing a standard Docker API where if any tool which communicates with the Docker daemon can use Docker Swarm to transparently scale to multiple hosts: Dokku, Docker Compose, Krane, Flynn, Deis, DockerUI, Shipyard, Drone, Jenkins and, of course, the Docker client itself. Docker Networking, Volumes and plugins can also be used through their respective Docker commands via Swarm. Swarm has been tested and is production ready to scale up to one thousand (1,000) nodes and fifty thousand (50,000) containers with no performance degradation in spinning up incremental containers onto the node cluster. Swarm also comes with a built-in scheduler, but you can easily plugin the Mesos or Kubernetes backend while still using the Docker client for a consistent developer experience. To find nodes in your cluster, Docker Swarm can use either a hosted discovery service, static file, etcd, consul and zookeeper depending on what is best suited for your environment.

ECOSYSTEM SUPPORT TO DOCKER

[3] Finally, it's easy to deploy Docker containers in a cloud scenario. You can easily integrate it with typical DevOps environments seamlessly (Ansible, Puppet, etc.) or use it as a standalone. The main reason it's so popular is simplification, says Ben Lloyd Pearson via opensource.com. You can do local development within a system that is identical to a live server; deploy various development environments from your host that each use their own software, OS, and settings; easily run tests on various servers; and create an identical set of configurations, so that collaborative work isn't ever hindered by parameters of the local host. Ecosystem support for Docker is improving with every passing day as it is gaining the popularity among the developers and the system operators.

Operating Systems support to Docker: Is compatible with virtually any distribution with a 2.6.32 + kernel RedHat Docker collaboration to work with across Fedora and other (2.6.32)+

It is compatible with Private PaaS (Platform as a Service) technologies: OpenShift, Solum (Rackspace and OpenStack).

Public PaaS technologies like Voxoz, Cocaine (Yandex), Biadu, etc.

Ecosystem Support for IaaS is present via Rackspace, Digital Ocean, AWS (Amazon Web Services), AML, etc.

Orchestration tools Support and Integration with: Chef, Puppet, Jenkins, Travis, Ansible..

In OpenStack Docker integration into NOVA (compatibility with Glance, Horizon) are also present

SHORTFALLS OF DOCKER

Though Docker is a legacy virtualization technique, they are not a go-to solution for all kinds of virtualization and cannot be considered as a replacement of the VMs.

VMs are self-contained with as they have a unique operating system (OS), drivers and application components whereas the containers are dependent as containers under Docker cannot run on Windows server.

VMs provide a high level of isolation as the system's underlying hardware resources are all virtualized and thus any bugs or viruses could not affect the other VM in the virtualized environment.

The kinds of tools and technologies required to manage the containers are still lacking in the industry and thus only a few management tools from companies like Google and Docker like Kubernetes and Swarm respectively are present which is not a good situation for an open-source product.

FUTURE WITH DOCKER

Docker Inc has set a clear path on the development of core capabilities (libcontainer), cross service management (libswarm) and messaging between containers (libchan). Docker's key open source project is libcontainers, which is on its way to becoming the default standard for Linux-based technology. Libcontainers enables the containers to work with the Linux namespace, control groups, capabilities, AppArmor security profiles, network interfaces and firewalling rules in a consistent and predictable way. [3] Libcontainers is getting help from Google, RedHat, and Parallels to build the program as they will work with Docker as core maintainers of the code. libcontainer which is written natively in Google's Go, is also being ported into other languages. Microsoft may be porting it to ASP.NET. Parallels' libct, which includes libcontainer's functionality, has native C/C++

and Python bindings. Microsoft is even jumping on board by bringing Docker to their Azure platform, a development that could potentially make integration of Linux applications with Microsoft products easier than ever before.

CONCLUSION

[6] With more than 1200 Docker Contributors, 10,000 Dockerized Applications at index.docker.io, 3 to 4 million Developers using Docker, 300 Million Downloads, 32,000 Docker related Projects and 70 percent of enterprise using Docker! It is not an exaggeration that with Docker Ecosystem there is a greater potential for some advanced deployment tools that combine containers, configuration management, continuous integration, continuous delivery and service orchestration in the days ahead.

REFERENCES

- [1] CenturyLink, "What is docker and when to use it," Web Page, Apr. 2014, accessed 2017-2-20. [Online]. Available: <https://www.clt.io/developers/blog/post/what-is-docker-and-when-to-use-it/>
- [2] Ken Cochrane, "How is docker different from a normal virtual machine?" Web page, apr 2013, accessed 2017-02-23. [Online]. Available: <http://stackoverflow.com/questions/16047306/how-is-docker-different-from-a-normal-virtual-machine>
- [3] opensource.com, "Who's using docker?" Web Page, Jul. 2014, accessed 2017-2-19. [Online]. Available: <https://opensource.com/business/14/7/docker-through-hype>
- [4] Docker Inc, "What is docker machine," Web Page, accessed 2017-2-23. [Online]. Available: <https://docs.docker.com/machine/overview/>
- [5] —, "Docker swarm," Web Page, accessed 2017-2-21. [Online]. Available: <https://www.docker.com/products/docker-swarm>
- [6] DATADOG, "8 surprising facts about real docker adoption," Web Page, accessed 2017-2-23. [Online]. Available: <https://www.datadoghq.com/docker-adoption/>