Hadoop YARN

Milind Suryawanshi[1,2] **and Gregor von Laszewski**[1,*]

[1] *School of Informatics and Computing, Bloomington, IN 47408, U.S.A.*
[2] *Savitribai Phule Pune University 2010, Pune, Maharashtra 411007 India*
[*] *Corresponding authors: laszewski@gmail.com*

Apache Hadoop 2.0 came up with Yarn architecture which is capable of managing resources and processing of tasks separately. This is the most important implementation over MapReduce which increased the scalability of Hadoop 1.0. YARN is the next generation of Hadoop's compute platform.

Keywords: Hadoop, Yarn, MapReduce, Cluster

https://github.com/MilindSuryawanshi/sp17-i524/tree/master/paper2/S17-IO-3020/report.pdf

---

> This review document is provided for you to achieve your best. We have listed a number of obvious opportunities for improvement. When improving it, please keep this copy untouched and instead focus on improving report.tex. The review does not include all possible improvement suggestions and if you sea comment you may want to check if this comment applies elsewhere in the document.

> I am surprised that your paper still have compiling error. This is already second paper and you can't do your reference right. The year field within each reference is wrong and causing error but you just ignore the message?

> project 000 should be replaced with paper2.

> Please apply 80 character width in your editor.

1. INTRODUCTION

Apache Hadoop

> Citation

is open source software framework, can be installed on a cluster of computers, which can communicate together for large amount of data storage and processing it, in highly distributed manner. Hadoop 2.0 project was developed to resolve the limitations of MapReduce (we will see the limitation as we move ahead). YARN (Yet Another Resource Negotiator) is Apache Hadoop's cluster resource management system [1]. It's a resource management technology which makes pace between, the way applications use Hadoop system resources and node manager agents. To understand Yarn, it is important to know the limitation of Hadoop1.0 MapReduce, which enforced the creation of Hadoop 2.0

> Missing period.

2. LIMITATION OF MAPREDUCE

Hadoop 1.0 was designed for big data processing and MapReduce was the only option supported. In it, Node holds the meta data information and the daemon Job Tracker will be ensuring that the Task Tracker are processing the data. Task Tracker runs the Map

> Should explain what is Map.

and reduces the tasks and reports the status to Job Tracker. Job Tracker tracks the tasks, schedules the jobs and monitors the Tasks Tracker. As you clearly see,

> We don't talk like this in a paper.

Job Tracker performs multiple operations which is bottleneck as we keep increasing the number of Task Trackers. Maximum 4000 Task Tracker and 40,000 concurrent tasks, can be handled by single Job Tracker. This was causing the limitation on use of number of Tasks Trackers. In MapReduce, there was a hard partition on resource utilization, which was causing inefficient use of resources [2].

> The figure title is inappropriate.

3. DESIGN

To overcome the limitation of MapReduce, Hadoop 2.0 was introduced with Yarn. Yarn is an architectural center of Hadoop 2.0 design [3]. It allows multiple data processing engines like Batch Streaming

> Citation

, Interactive SQL

> Citation

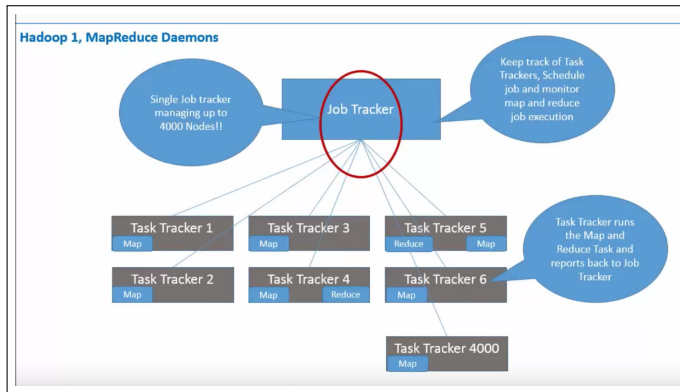, Data science and Real time streaming to store and retrieve

Fig. 1. Source [2].

data in Hadoop Distributed File System (HDFS). Yarn is prerequisite to Enterprise Hadoop. It acts as middle layer in Hadoop cluster which extends the Hadoop capability which provides resource management, delivers consistent operation, security and data governance tools across Hadoop cluster. It provides ISV

> What is the full name for ISV?

engines with a consistent framework which allows developer to write data access applications that run in Hadoop.
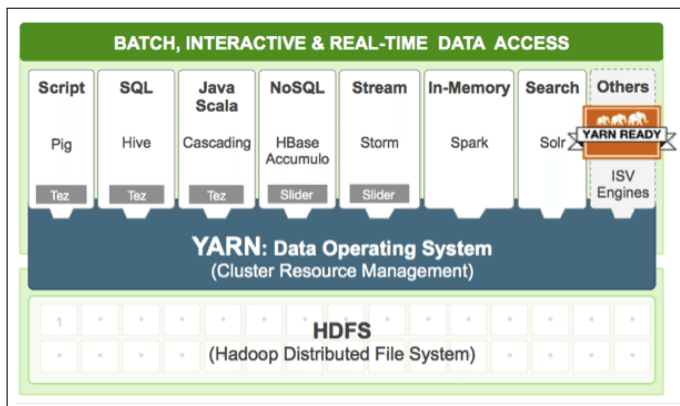

Fig. 2. Source [3].

> Once again, the figure title shouldn't be "source". Also, when you place a figure within your paper without introducing it in your paper, why do you put it here?

4. ARCHITECTURE

Apache Hadoop YARN architecture has ResourceManager, ApplicationManager, NodeManager and distributed application as its important parts. YARN's global ResourceManager runs in a master daemon. Usually ResourceManager is placed in dedicated machine, it allocates the requested resource based on available cluster resources. It keeps track of how many live nodes and resources are available on cluster for allocation. It coordinates with the NodeManager for resource allocation and release functionality. ResourceManager consists of two major components Scheduler and ApplicationManager [4]. Scheduler is pure scheduler which only allocates the resource, it will not perform any monitoring or tracking of status. Scheduler schedules the

resource (like cpu, disk, network etc.) based on applications requirement. Scheduler can divide the cluster resource among the application tasks, queues it is called as pluggable policy [4]. Once user submits the application to the ResourceManager, the ResourceManager's scheduler will allocate the resource to it. Once resource get allocated ApplicationManger comes into picture and perform the coordination all the tasks running for that application like: monitoring of tasks, restarting failed tasks, speculatively running slow tasks and total values of application counter. It also asks for appropriate resources to run tasks [3].
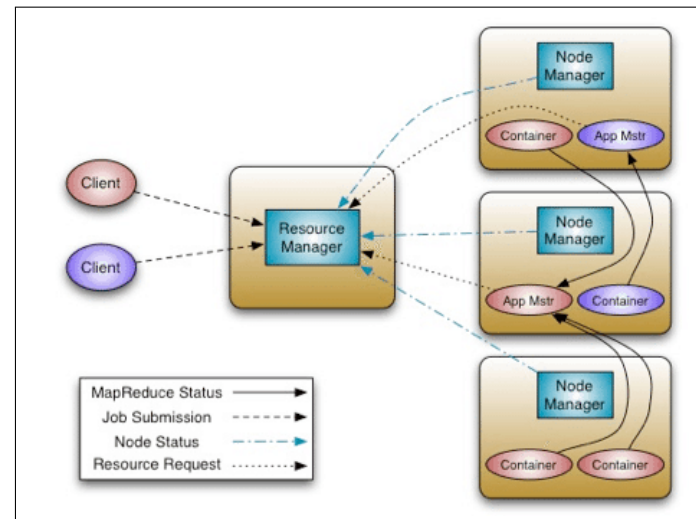

Fig. 3. Source [1].

> Another wrong figure title. Also, please describe the figure within the text.

NodeManager will create the container to perform the given tasks. Container size depends on the tasks he needs to perform. Container can be any resource type like CPU, disk, network, and storage. NodeManger can have number of containers depending on the configuration parameter and node resource capacity. Each application has its ApplicationManager. It performs required tasks inside a container. E.g. the MapReduce. ApplicationMaster performs the map and reduce tasks. On similar line Giraph's ApplicationMaster perform Giraph specific tasks in a container. The ResourceManger, the NodeManager and a container work regardless of the type of application. This Yarn feature, makes it more popular, as this allows to run different types of application in a single cluster. This generic approach allows Hadoop Yarn cluster to run various application like MapReduce, Giraph, Storm, Spark, Tez/Impala, MPI etc. To understand the popularity of the Hadoop Yarn, we need to know its advantages.

4.1. Advantages

- Resources can be shared among different clusters, which maximizes the efficiency of resource utilization.
- All in one cluster will help to run maximum application, will reduce the operation cost.
- Reduce the data transfer operation, as there is no need to transfer the data between Hadoop Yarn and system running on different clusters of machine [3].

> Generally, bullet point is not welcome in paper.

5. FEATURE

Yarn came with lots of features; we are not covering each of the features here. However, this article will be covering the most important feature, compared to MapReduce, which makes Yarn popular [3].

- **Uberization:** To reduce the overhead on ResourceManager, a small tasks container can ask NodeManager to start their tasks. So all the small tasks of MapReduce are run by ApplicationMaster's JVM. This improves the performance.

- **Multi-tenancy:** Yarn allows multiple engines to access the Hadoop as a common platform for batch processing, interactive and real time data access. This feature returns the most enterprise investment for business.

- **Cluster Utilization:** This feature overcomes the limitation of not using hardware resource efficiently for jobs in MapReduce. Yarn dynamically allocates the resource which improves the resource utilization.

- **Scalability:** Introduces cluster manager so ResourceManager can solely focus on scheduling and let cluster to track the live nodes and available resources in the cluster and assign the tasks to them.

- **Computability:** Existing MapReduce application can run on the latest Hadoop2 Yarn, without failing.

6. LIMITATION

Any new offering takes time to mature and align with market expectation. There are following limitations of Yarn. As the product is growing it is coming with better ecosystem to overcome its limitation [2]:

- Complexity

  - Yarn is complex and level of abstraction is low, from the perspective of developers.

  - Developers need to do very low level tasks, for example: a Hello World program is of 1500 lines code.

  - Clients need to be prepared with all the dependencies (it is mostly library files which help the current program to run with additional functionality).

  - To up and run the app, client requires to setup the environment, like class path.

  - To see the log, developer needs to wait till completion of the job, currently there is no mechanism to show the console log while the job is running.

- Application cannot handle the master crash, which causes the single point failure of the app

- There is no *built in* communication layer available between master and container.

- Not helpful in long running Jobs

- Hard to debug

7. CONCLUSION

Apache Hadoop Yarn is important part of Hadoop 2.0 architecture, which separates the resource management and processing components. With this functionality it achieves scalability, efficiency and flexibility compared to MapReduce engine (from first version of Hadoop). Adding to this Yarn based architecture will not limit the existing MapReduce applications. It supports such, existing application to run in Hadoop 2.0, without failure. Most of the Hadoop 1.0 users are migrating to Hadoop 2.0 Yarn and using it successfully, such as YAHOO!, eBay, Spotify, Xing, Allegro and more.

Use case on Yarn is missing.

REFERENCES

[1] Arun Murthy, "Apache Hadoop Yarn Background An Overview," Web Page, 2012, page last accessed on 7 April 2017. [Online]. Available: https://hortonworks.com/blog/apache-hadoop-yarn-background-and-an-overview/

[2] BigDSol, "Understand Hadoop YARN," YouTube, 2013, page last accessed on 7 April 2017. [Online]. Available: https://www.youtube.com/watch?v=1vg_W-MMZpA

[3] Horton works, "APACHE HADOOP YARN," Web Page, page last accessed on 7 April 2017. [Online]. Available: https://hortonworks.com/apache/yarn/

[4] The Apache Software Foundation, "Apache Hadoop Yarn," Web Page, 2016, page last accessed on 7 April 2017. [Online]. Available: https://hadoop.apache.org/docs/r2.7.2/hadoop-yarn/hadoop-yarn-site/YARN.html