

Heroku

YATIN SHARMA^{1,*}, +

¹ School of Informatics and Computing, Bloomington, IN 47408, U.S.A.

* Corresponding authors: yatins@indiana.edu

+ HID - S17-IR-2034

Paper-002, March 26, 2017

Heroku is a cloud application program that enables developers to build and deploy application in almost any language. It provides everything a developer would need to build customer facing application. The core of Heroku is called Dyno, which can be thought of as virtual machine that runs just your application. Dynos are scalable according to our needs as well.

© 2017 <https://creativecommons.org/licenses/>. The authors verify that the text is not plagiarized.

Keywords: Heroku, PaaS, Dyno, Dyno Manifold, Logplex, Toolbelt client, Procfile.

<https://github.com/yatinsharma7/sp17-i524/tree/master/paper2/S17-IR-2034/report.pdf>

INTRODUCTION

Heroku[1] is Platform as a Service(PaaS)[2] that enables us to build application in any language on demand. It allows us to deploy web applications seamlessly as well as monitor and share with other developers at the click of a button. It is all about rapid application development for the cloud, using the underlying platform infrastructure and software add-ons to build, deploy and monitor large and scalable web application. At present it handles 5 billion requests per day.

BENEFITS

Heroku is language agnostics and provides great flexibility in choosing an appropriate programming language to develop the web application. It has core support for Ruby, Java, Python, Node.js, and PHP. Also, any other language can be supported by using a feature called buildpacks. Heroku provides a lot of flexibility in managing our applications after deployment using the Heroku command line tool running on the client machine or on the Heroku Infrastructure. Heroku is Git[3] focused and makes it easy to share code using version controlling. Heroku is also a part of the Salesforce family. It has a feature called Heroku Connect that allows to build application that share data with the Salesforce deployment.

ARCHITECTURE

The Heroku architecture consists of platform stack containing various runtime libraries, OS, and underlying infrastructure. The unit of work in Heroku framework is called Dyno. It can be thought of as packaged running version of our code that the application interacts with. Dynos are responsible for receiving web requests, writing an output and connecting to application resources such as databases. They are fully isolated containers

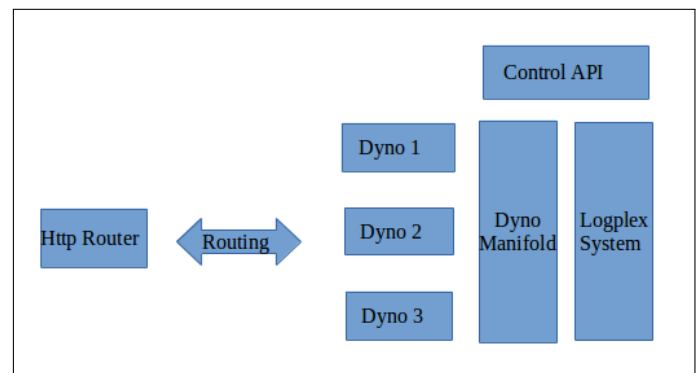


Fig. 1. High Level Architecture of Heroku platform

running on the Dyno Manifold, which is the building block for execution environment. Dyno Manifold is responsible for process management, isolation, scaling, routing, distribution that are necessary for to run the web and worker processes. It is also fault tolerant and distributed in nature. If any dynos fails, the manifold restarts them automatically, hence removing a lot of maintenance hassles.

Dyno Execution

Dynos are capable of serving many request per second and execute in complete isolation from each other. Each dyno gets its own virtual environment that it can use to handle its own client requests. Dynos use LXC to provide container like behavior to achieve complete isolation from one another. There is a memory restriction of 1.5 GB per dyno, beyond which the dyno is rebooted with a Heroku error which could lead to a memory leak.

Execution Flow

Process type is the declaration of the command that defines the structure to be used while instantiating a process. Heroku has two process types- web process, which is responsible for handling HTTP client requests and the worker process, which is responsible for executing other tasks such as customer jobs of running background jobs and queuing tasks.

Logging Architecture

Heroku logplex system provides a flexibility facility by giving us an overall view of our application runtime behavior. It forms the basis of the Heroku logging infrastructure. It routes log streams from various sources into single output (for example archival system). The logplex system keeps the most recent data (typically 1500 logs) that are important to extract relevant information from the application being run.

Http Routing

Routing Mesh are responsible for routing the web requests to the appropriate web process dyno. It is also responsible for seeking the application's web dynos within the dyno manifold and forwarding the HTTP web requests to one of them. The routing mesh uses a round robin algorithm to distribute the request across various dynos. Since the dynos could be running in distributed manner, the routing mesh manages the access and routing internally and none of the dynos are statically addressable. Heroku also supports multithreaded and asynchronous application, accepting many network connections to process client requests.

PROCESS ARCHITECTURE

A Heroku application can be thought of a multiple process, each consuming resources like a normal UNIX process, that run on the Heroku Dyno manifold. Heroku defines each process through configuration file called Procfile- which is a text file, placed in the root of the application and contains the format describing how our application will run.

HEROKU PLATFORM API

Heroku platform API is a tool that enables us to call Heroku platform services and create application, plug in new add-ons simply using HTTP. It gives the developers complete control over their application. The three components that define the behavior of the API are : 1)Security 2)Schema 3)Data. The client accesses the API using standard methods defined for HTTP. The API then acts on the request and returns the result in JSON format.

SECURITY

Heroku employs various measures to ensure that the application and data stores within the platform are secure from external attacks, thefts and hacks. Heroku enforces SSH[4] protocol to encrypt the source code while they are getting pushed into the Heroku environment. Any application that runs on Heroku is in complete isolation from each other, so that no two application can see each other getting executed. It also restricts applications from making local network connection between hosts. It enables data security by keeping the data in access controlled databases.

GETTING STARTED

There are few prerequisites that we need to perform before we can start using Heroku: 1) Get Heroku account 2) Install Heroku toolbelt client[5] 3) Set up SSH for our user account. Heroku toolbelt is the client software required to work with the Heroku platform and contains the following component: Heroku Client, Foreman and Git.

CONCLUSION

Heroku is dramatically different from the traditional hosting or any normal cloud infrastructure offerings. Instead of thinking about virtual servers as individual units and trying to figure out how many do we need and the communication between them, Heroku completely abstracts the servers and filesystems away. As a developer we just have to push the code and heroku will manage all the rest of the process. Heroku provides a complete developer experience and application runtime and also frees the developer from hassles of underlying infrastructure. It manages all that in scalable and highly maintainable fashion.

ACKNOWLEDGEMENTS

The author thanks Prof. Gregor von Laszewski for his technical guidance.

REFERENCES

- [1] "Cloud application platform | heroku," Web Page, online; accessed 13-Mar-2017. [Online]. Available: <https://www.heroku.com/>
- [2] "Platform as a service - wikipedia," Web Page, online; accessed 13-Mar-2017. [Online]. Available: https://en.wikipedia.org/wiki/Platform_as_a_service
- [3] "Github," Web Page, online; accessed 13-Mar-2017. [Online]. Available: <https://github.com/>
- [4] "Secure shell- wikipedia," Web Page, online; accessed 17-Mar-2017. [Online]. Available: https://en.wikipedia.org/wiki/Secure_Shell
- [5] "Heroku cli | heroku dev center," Web Page, online; accessed 17-Mar-2017. [Online]. Available: <https://devcenter.heroku.com/articles/heroku-cli>