

Heroku

YATIN SHARMA^{1,*}, +

¹ School of Informatics and Computing, Bloomington, IN 47408, U.S.A.

* Corresponding authors: yatins@indiana.edu

+ HID - S17-IR-2034

Paper-002, April 12, 2017

Heroku is a web application hosting platform-as-a-service cloud that enables developers to build and deploy application. It supports multiple programming languages including Ruby, Java, Node .js and Python. It is a simple and modular platform that allows developers to focus less on infrastructure/deployment and focus more on coding.

© 2017 <https://creativecommons.org/licenses/>. The authors verify that the text is not plagiarized.

Keywords: Heroku, PaaS, Dyno, Dyno Manifold, Logplex, Toolbelt client, Procfile.

<https://github.com/cloudmesh/sp17-i524/tree/master/paper2/S17-IR-2034/report.pdf>

1. INTRODUCTION

Heroku[1] is Platform as a Service(PaaS)[2] that enables us to build and deploy applications in various programming language including Java, Python, Ruby, etc. It allows us to deploy web applications seamlessly as well as monitor and share with other developers instantly. It is enables rapid application development for the cloud, using the underlying platform infrastructure and software add-ons to build, deploy and monitor large and scalable web applications. It is built on top of Amazon Web Services and is owned by Salesforce.com [3].

2. BENEFITS

Heroku is language agnostics and provides great flexibility in choosing an appropriate programming language to develop the web application. It has core support for Clojure, Ruby, Java, Python, Node.js, Scala and Play, Go, and PHP. Apart from these, any other language can also be supported by using a feature called buildpacks. Heroku provides a lot of flexibility in managing the applications after deployment using the Heroku command line tool running on the client machine or on the Heroku Infrastructure. Heroku uses Git[4] as the means for deploying applications to its servers. Being owned by Salesforce.com[3], it has a feature called Heroku Connect enabling interaction between the applications and Salesforce API.

3. ARCHITECTURE

The Heroku architecture consists of platform stack containing various runtime libraries, OS, and underlying infrastructure. A Heroku application can be thought of a multiple processes, each consuming resources like a normal UNIX process, that run on the Heroku Dyno manifold. Heroku defines each process through configuration file called Procfile- which is a text file, placed in

root of the application and contains the format describing how the application will run. Fig. 1. below describes the high level architecture of the Heroku platform.

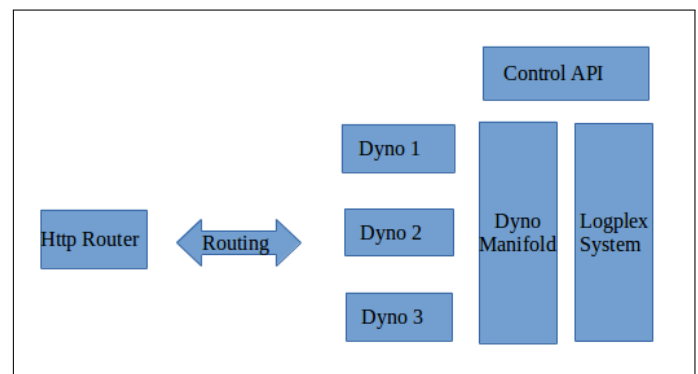


Fig. 1. High Level Architecture of Heroku platform

3.1. Process Management

The unit of work in Heroku framework is called Dyno. It can be thought of as packaged running version of the code that the application interacts with. Dynos are responsible for receiving web requests, writing an output and connecting to application resources such as databases. They are fully isolated containers running on the Dyno Manifold, which is the building block for execution environment. Dyno Manifold is responsible for process management, isolation, scaling, routing, distribution that are necessary for to run the web and worker processes. It is also fault tolerant and distributed in nature. If any dynos fails, the manifold restarts them automatically, hence removing a lot of maintenance hassles. Dynos are capable of serving many

request per second and execute in complete isolation from each other. Each dyno gets its own virtual environment that it can use to handle its own client requests. Dynos use LXC to provide container like behavior to achieve complete isolation from one another. There is a memory restriction of 1.5 GB per dyno, beyond which the dyno is rebooted with a Heroku error which could lead to a memory leak.

3.2. Execution Flow

Process type is the declaration of the command that defines the structure to be used while instantiating a process. Heroku has two process types- web process, which is responsible for handling HTTP client requests and the worker process, which is responsible for executing other tasks such as customer jobs of running background jobs and queuing tasks.

3.3. Logging Architecture

Heroku logplex system provides a flexibility facility by giving us an overall view of the application runtime behavior. It forms the basis of the Heroku logging infrastructure. It routes log streams from various sources into single output (for example archival system). The logplex system keeps the most recent data (typically 1500 logs) that are important to extract relevant information from the application being run.

3.4. Http Routing

Routing Mesh are responsible for routing the web requests to the appropriate web process dyno. It is also responsible for seeking the application's web dynos within the dyno manifold and forwarding the HTTP web requests to one of them. The routing mesh uses a round robin algorithm to distribute the request across various dynos. Since the dynos could be running in distributed manner, the routing mesh manages the access and routing internally and none of the dynos are statically addressable. Heroku also supports multithreaded and asynchronous application, accepting many network connections to process client requests.

3.5. Heroku Interfaces

Heroku provides the developer with the flexibility to control various aspects of the application through various control surfaces such as process management, routing, logging, scaling, configuration, and deployment. These are available as a command-line interface (CLI), a web-based console, and full REST API.

4. ADDONS

Heroku supports vast amount of third party addons[5] that any Heroku user can instantly provision as an attachable resource for their application. For example, if the application needs a PostgreSQL[6] database, it can be done in Heroku. Addons work through Heroku's environment variables. Each time an addon is added to the application, the application will automatically be assigned new environment variables which specify any configurations required to interact with the new addon. Just like Dynos, Addons are easy to add, up-grade, downgrade, and remove without requiring any downtime for the applications.

5. HEROKU PLATFORM API

Heroku platform API is a tool that allows to call Heroku platform services, create applications, and plug in new add-ons by simply using HTTP. It gives the developers complete control over their

application. The three components that define the behavior of the API are : 1)Security 2)Schema 3)Data. The client accesses the API using standard methods defined for HTTP. The API then acts on the request and returns the result in JSON format.

6. SECURITY

Heroku employs various measures to ensure that the application and data stores within the platform are secure from external attacks, thefts and hacks. Heroku enforces SSH[7] protocol to encrypt the source code while they are getting pushed into the Heroku environment. Any application that runs on Heroku is in complete isolation from another, so that no two application can see each other getting executed. It also restricts applications from making local network connection between hosts. It enables data security by keeping the data in access controlled databases.

7. GETTING STARTED

There are few prerequisites that we need to perform before we can start using Heroku: 1) Get Heroku account 2) Install Heroku toolbelt client[8] 3) Set up SSH for the user account. Heroku toolbelt is the client software required to work with the Heroku platform and contains the following component: Heroku Client, Foreman and Git. Step-by-step procedures to download, install and get started with Heroku can be found on-line.[9]

8. CONCLUSION

Heroku offers a Platform as a Service (PaaS) in which the servers and filesystems are completely abstracted away. It gives provides an environment where pushing the code and basic configuration can get an application running. It provides a complete developer experience and an application runtime. It manages all that in scalable and highly maintainable fashion. It's logging service as well as instant scaling power allows scaling up to support demand from users, and scaling down when high traffic has stopped.

ACKNOWLEDGEMENTS

The author thanks Prof. Gregor von Laszewski for his technical guidance.

REFERENCES

- [1] "Cloud application platform | heroku," Web Page, online; accessed 13-Mar-2017. [Online]. Available: <https://www.heroku.com/>
- [2] "Platform as a service - wikipedia," Web Page, online; accessed 13-Mar-2017. [Online]. Available: https://en.wikipedia.org/wiki/Platform_as_a_service
- [3] "Salesforce.com | the customer success platform to grow your business," Web Page, online; accessed 3-April-2017. [Online]. Available: <https://www.salesforce.com/>
- [4] "Github," Web Page, online; accessed 13-Mar-2017. [Online]. Available: <https://github.com/>
- [5] "Add-ons - heroku elements," Web Page, online; accessed 3-April-2017. [Online]. Available: <https://elements.heroku.com/addons>
- [6] "Heroku postgres - add-ons - heroku elements," Web Page, online; accessed 3-April-2017. [Online]. Available: <https://elements.heroku.com/addons/heroku-postgresql>
- [7] "Secure shell- wikipedia," Web Page, online; accessed 17-Mar-2017. [Online]. Available: https://en.wikipedia.org/wiki/Secure_Shell
- [8] "Heroku cli | heroku dev center," Web Page, online; accessed 17-Mar-2017. [Online]. Available: <https://devcenter.heroku.com/articles/heroku-cli>

- [9] "Getting started on heroku | heroku dev center," Web Page, online; accessed 3-April-2017. [Online]. Available: <https://devcenter.heroku.com/start>